

Chapter 2: Multi-armed Bandits - Highlights

Jacob Taylor Cassidy

November 22, 2022

1 A k-armed Bandit Problem

- .

2 Action-value Methods

- .

3 The 10-armed Testbed

- .

4 Incremental Implementation

- .

5 Tracking a Nonstationary Problem

- .

6 Optimistic Initial Values

- .

7 Upper-Confidence-Bound Action Selection

- .

8 Gradient Bandit Algorithms

- .

9 Associative Search (Contextual Bandits)

- .

10 Summary

- .